I implemented the naive bayes classifier pretty much straight from the book, with the optimizations suggested for sentiment analysis, specifically the binary word counts and prepending ‘NOT_’ to all words after a negation. I only used standard python libraries, re for regex to parse the text, Counter from collections to quickly and easily create a dictionary of word counts, operator to provide a one line solution to get argmax_c(sum[c]), and math for log(), string to get python’s preconstructed list of punctuation.